

LangChain, LlamaIndex and CrewAI

What, How, When ?

All three are LLM orchestration frameworks. They connect models to your data, tools and workflows. They overlap, and choosing between them is easier once you see how each one runs.

[swipe →](#)

Each one is built around something different

- **LangChain**

built around: steps & state

A toolkit for wiring models, tools and memory together. LangGraph, its lower-level runtime, runs agents as a graph you control.

- **LlamaIndex**

built around: documents & queries


A data framework for getting your own documents into an LLM: load, parse, index, retrieve, answer with sources.

- **CrewAI**

built around: roles & tasks


A multi-agent framework where you describe a team, give them tasks, and let the crew work through them.

An agent is a loop in a graph




- The model decides whether to call a tool or answer.
- Tool results always go back to the model. Only the model can end the loop.
- Optional human review pauses before a tool runs: approve, edit or reject.

Ingest once, retrieve on every question



- Load your data, parse messy files, then chunk, embed and index it.
- For each question, the retriever finds the most relevant chunks.
- The model answers from those chunks and can cite them as sources.

Describe a team, give it tasks



- Each agent has a role, a goal and a backstory.
- Each task says what output is expected, and can use earlier tasks as context.
- Tasks run in order, or a manager agent delegates them (hierarchical mode).

Which one should you reach for ?

Chat with your company's PDFs and docs

- **LlamaIndex**

An agent with branches, retries and human approval

- **LangChain / LangGraph**

A team of agents: researcher, writer, reviewer

- **CrewAI**

Extract fields from messy invoices and forms

- **LlamaIndex**

A quick multi-step demo for tomorrow

- **CrewAI**

A production agent you need to debug and monitor

- **LangChain / LangGraph**

In real projects people sometimes mix them. A LlamaIndex retriever as a tool inside a LangGraph or CrewAI agent is a common setup.

The differences at a glance

CORE ABSTRACTION		
LangChain Agent loop built on a graph of nodes, edges and shared state	LlamaIndex Index, retriever and query engine, plus event-driven Workflows	CrewAI Agents with a role, goal and backstory, Tasks, and a Crew
CONTROL FLOW		
LangChain You define it explicitly	LlamaIndex Mostly a pipeline. Workflows add steps and events	CrewAI Sequential, hierarchical, or Flows for explicit control
STRONGEST AT		
LangChain Custom, stateful agents with full control	LlamaIndex RAG and document-heavy apps	CrewAI Getting a team of role-based agents working fast
TRADE-OFF		
LangChain More concepts to learn, and you design the graph yourself	LlamaIndex Less focused on multi-agent orchestration	CrewAI Less step-by-step control unless you use Flows

What does it cost ?

FRAMEWORK COST		
LangChain	LlamaIndex	CrewAI
Free, open source (MIT)	Free, open source (MIT)	Free, open source (MIT)
WHAT YOU PAY FOR		
LangChain	LlamaIndex	CrewAI
LLM calls. One per pass through the loop, so an agent using 3 tools makes at least 4 calls	LLM calls. Embeddings once when you index, then about one call per question	LLM calls. At least one run per agent, plus the manager's calls in hierarchical mode
PAID PLATFORM (AS OF SEPT 2026)		
LangChain	LlamaIndex	CrewAI
LangSmith: free for 1 seat and 5k traces a month. Plus is \$39 per seat a month	LlamaParse: 10k free credits a month. Starter \$50 a month, Pro \$500 a month. 1,000 credits = \$1.25	CrewAI platform: 50 free workflow runs a month. Enterprise has custom pricing

The frameworks are free. Your bill depends mostly on the model you choose and how many calls your design makes.

REFERENCES

Where to read more

PAPERS

1. Rasheed et al. (2026). Agentic Frameworks for Reasoning Tasks: An Empirical Study.
arxiv.org/abs/2604.16646
2. Ferrag, Tihanyi and Debbah (2025). From LLM Reasoning to Autonomous AI Agents: A Comprehensive Review.
arxiv.org/abs/2504.19678
3. Singh et al. (2025). Agentic Retrieval-Augmented Generation: A Survey on Agentic RAG.
arxiv.org/abs/2501.09136

DOCUMENTATION

1. LangChain docs.
docs.langchain.com
2. LangChain and LangGraph 1.0 release.
langchain.com/blog/langchain-langgraph-1dot0
3. LlamaIndex docs.
developers.llamaindex.ai
4. CrewAI docs.
docs.crewai.com

PRICING

1. LangSmith pricing.
langchain.com/pricing
2. LlamaParse and LlamaCloud pricing.
llamaindex.ai/pricing
3. CrewAI pricing.
crewai.com/pricing