

SWIPE TO LEARN

Vector Databases Explained

The complete guide: embeddings, similarity search, indexing algorithms, RAG, scaling, and the tools powering modern AI — from first principles.



Normal Databases Can't Search by **Meaning**

A regular database finds **exact matches**. Search "running shoes" and it will miss a product called "jogging sneakers" — even though they mean the same thing.

Humans don't think in exact keywords. We think in **meaning and context**. AI systems need a way to search the same way humans think.

This is exactly the gap that **vector databases** were built to close.

What Is a Vector Database?


A vector database stores data as **lists of numbers called vectors** — and lets you search by how similar two pieces of data are in meaning, not by exact text.

Simple analogy: Imagine every book in a library placed on a shelf based on its *topic*, not its title. Books about "space travel" sit near books about "rockets" — even if neither word appears in the other's title.

That is exactly what a vector database does — but for any kind of data: text, images, audio, or code.

What Is an Embedding?

An embedding is a piece of data turned into a **list of numbers** — a vector — that captures its meaning. An AI model creates this list.



Notice the numbers are **close to each other** — because the meanings are close. That closeness is what makes search possible.

How Embedding Models Are Trained

An embedding model is a neural network trained on huge amounts of data. During training, it learns to place **similar things close together** and different things far apart in number-space.

It does this by seeing millions of examples — pairs of sentences that mean the same thing, images that show the same object, or words that appear in similar contexts — and adjusting itself until "closeness" reflects real meaning.

The result: a reusable function that turns any new piece of data into a meaningful vector in a fraction of a second.

What Are Vector Dimensions?

Each number in a vector is called a **dimension**. A typical embedding has anywhere from 128 to 3,072 dimensions — each one capturing a subtle aspect of meaning.

1,536 dimensions

=

A point in 1,536-dimensional
space

More dimensions can capture more nuance, but they also cost more memory and slow down search — which is why choosing the right embedding size matters.

Finding the Nearest Neighbors

Once everything is a vector, search becomes a **math problem**: find the vectors that are closest to your query vector in space.

This is called Approximate Nearest Neighbor (ANN) search. Instead of checking every single item one by one — which is slow — the database uses smart shortcuts to find close matches almost instantly, even across billions of vectors.

The most common way to measure "closeness" is called **cosine similarity** — it measures the angle between two vectors. Smaller angle, more similar meaning.

Three Ways to Measure "Closeness"

1

Cosine Similarity

Measures the angle between vectors. Best when direction matters more than size.

2

Euclidean Distance

Measures the straight-line distance between two points in space.

3

Dot Product

Multiplies and sums vector values — fast, and factors in vector magnitude too.

Cosine Similarity

Explained Simply

Cosine similarity ignores **how long** a vector is and only looks at the **direction** it points. Two vectors pointing the same way score close to 1 — completely unrelated ones score close to 0.

Same direction



Score ≈ 1.0

It's the default metric for most text embeddings because it's not thrown off by how long or short a piece of text is.

Euclidean Distance

The Straight Line

Euclidean distance is the metric you learned in school — the **straight-line distance** between two points. The smaller the distance, the more similar the two vectors.

It works well for embeddings where actual position in space matters, like image features or spatial data — but it can be sensitive to vector length.

Most vector databases let you pick whichever metric fits your embedding model, since different models are trained to optimize for different ones.

Dot Product Similarity

The dot product multiplies matching dimensions of two vectors and adds up the results. Unlike cosine similarity, it's affected by **both direction and magnitude**.

It's computationally the **cheapest** metric to calculate, which is why many large-scale systems default to it for speed — especially when embeddings are already normalized.

Choosing the right metric can meaningfully change your search quality, so it's worth matching it to how your embedding model was trained.

Vector DB vs Traditional DB

Traditional DB	Vector DB
Exact match search	Meaning-based search
Rows and columns	Vectors (lists of numbers)
SQL queries	Similarity queries
Good for structured data	Good for text, images, audio
Misses synonyms	Understands synonyms naturally

How It Works

Step by Step

1

Convert Data to Vectors

An embedding model turns your text, image, or audio into a vector.

2

Store in the Vector Index

The vector is saved in a structure built for fast similarity search.

3

Convert the Query Too

Your search query is also turned into a vector the same way.

4

Find the Closest Matches

The database returns the stored vectors most similar to your query.

The Data Ingestion Pipeline

1

Load

Pull in raw documents, PDFs, images, or database records.

2

Chunk

Break long content into smaller, meaningful pieces.

3

Embed

Convert each chunk into a vector using an embedding model.

4

Index

Store vectors plus metadata in a structure built for fast retrieval.

Why Chunking Strategy Matters

Fixed-Size Chunking

Splits text every N characters or tokens — simple but can cut sentences mid-thought.

Sentence-Based

Splits along sentence boundaries to keep ideas intact.

Semantic Chunking

Uses meaning shifts to decide where one topic ends and another begins.

Sliding Window

Overlaps chunks slightly so context isn't lost at the edges.

Choosing the Right Chunk Size

Too small, and chunks lose context — a sentence pulled out of its paragraph can be meaningless on its own.

Too large, and chunks blend multiple ideas together, making the embedding vague and search results less precise.

Most teams start around 200–500 tokens per chunk with some overlap, then tune based on real search quality — there's no universal "right" number.

Popular Embedding Models Compared

Model	Best For
OpenAI text-embedding-3	General-purpose, strong out of the box
Cohere Embed	Multilingual search
BERT / Sentence-BERT	Open source, self-hosted projects
CLIP	Matching text with images
Voyage AI	Domain-specific fine-tuning

What Is a Vector Index?

A vector index is a specialized data structure that organizes vectors so the database doesn't have to compare your query against every single one to find matches.

Think of it like the index at the back of a book — instead of reading every page, you jump straight to the relevant section.

Different index types trade off **speed, accuracy, and memory** in different ways — choosing one is one of the most important decisions in building a vector search system.

Brute Force

(Flat Index)

The simplest approach: compare the query vector against **every single stored vector**, one by one, and return the closest ones.

It's perfectly accurate — the true nearest neighbors, guaranteed — but it gets slow fast as the number of vectors grows into the millions.

Flat indexes are still useful for small datasets or as a baseline to measure other indexes' accuracy against.

Approximate Nearest Neighbor (ANN) Search

ANN algorithms sacrifice a tiny bit of accuracy for a **massive speed boost** — instead of checking everything, they use clever shortcuts to skip most of the search space.


This is what lets modern vector databases search across billions of vectors in milliseconds.

The three most popular ANN families are **HNSW, IVF, and LSH** — each with different trade-offs, covered next.

HNSW

Explained Simply

HNSW (Hierarchical Navigable Small World) builds a **layered graph** connecting similar vectors together — like a multi-level map with highways on top and local roads below.



Search "hops" from the coarse top layer down to fine detail, landing near the true answer in just a few steps. It's the most widely used index today because of its excellent speed-to-accuracy balance.

IVF

(Inverted File Index)

IVF groups vectors into **clusters** based on similarity, then stores which cluster each vector belongs to.

When a query comes in, the database only searches the handful of clusters closest to the query — instead of scanning everything — dramatically cutting down the work.

IVF is memory-efficient and scales well, which is why it's often paired with quantization for very large datasets.

Product

Quantization (PQ)

Quantization compresses vectors by rounding their values into a much smaller set of representative "codes" — trading a little precision for a huge cut in memory usage.

This lets you store **billions of vectors** in memory that would otherwise be far too large, at the cost of a small accuracy drop.

It's almost always used alongside another index like IVF or HNSW, rather than on its own.

LSH

(Locality Sensitive Hashing)

LSH uses special hash functions designed so that **similar vectors are more likely to land in the same bucket** — the opposite goal of a normal hash function, which tries to avoid collisions.

To search, you only need to check the items in the same bucket as your query, instead of the whole dataset.

LSH is simple and fast to build, though it's largely been overtaken by HNSW for most modern production use cases.


Index Trade-offs

At a Glance

Index	Best Trait
Flat / Brute Force	Perfect accuracy, small datasets
HNSW	Best speed-to-accuracy balance
IVF	Efficient at very large scale
PQ (Quantization)	Lowest memory footprint
LSH	Simple, fast to build

Metadata Filtering

Vectors alone can't express things like "only show results from the last 30 days" or "only from category: electronics." That's where metadata comes in.



Every stored vector can carry extra fields alongside it, so queries can combine **meaning-based search** with **exact filters** in a single step.

What Is Hybrid Search?

Hybrid search combines **keyword search** (great for exact terms, product codes, names) with **vector search** (great for meaning and intent).

Results from both methods are blended using a scoring formula, giving you the precision of keyword matching and the flexibility of semantic understanding at the same time.

Most production search systems today use hybrid search rather than vectors alone, because pure semantic search can miss exact codes or rare terms.

What Is Re-ranking?

Vector search is fast because it uses approximate shortcuts — but that means the very top results aren't always perfectly ordered.

Re-ranking takes the top handful of candidates returned by the vector search and runs a slower, more precise model over just those few to reorder them by true relevance.

It's a common final step in production RAG and search pipelines — fast recall first, precise ranking second.

RAG:

Retrieval-Augmented Generation

Large language models don't know your private documents, and they can "hallucinate" facts. RAG fixes this by **retrieving relevant information first**, then feeding it to the model as context.

The vector database is the retrieval engine: it finds the most relevant chunks of your data for a given question, in real time.

This is why nearly every modern AI chatbot that "knows your documents" is powered by a vector database behind the scenes.

The RAG Pipeline, Step by Step

1

User Asks a Question

The question is converted into a vector.

2

Retrieve Relevant Chunks

The vector database returns the most similar stored chunks.

3

Augment the Prompt

Those chunks are inserted into the prompt sent to the LLM.

4

Generate the Answer

The LLM answers using the retrieved facts as grounding.

Multi-Tenancy in Vector Databases

When you build a product for many customers, each customer's data usually needs to stay **completely isolated** — one user should never see another user's search results.

Vector databases handle this through **namespaces, partitions, or per-tenant collections**, so searches only ever run within one customer's data at a time.

Getting this right is essential for SaaS products handling private documents at scale.

Real-Time Updates and Freshness

Data changes constantly — new documents, edited products, deleted records. A good vector database supports **inserting, updating, and deleting** vectors without rebuilding the whole index.

Some ANN indexes handle live updates gracefully; others need periodic rebuilds to stay accurate and fast.

For fast-moving data — like live chat or news — freshness matters just as much as search quality.

Scaling With Sharding & Replication

Sharding splits your vectors across multiple machines so no single server needs to hold everything — essential once you pass millions or billions of vectors.

Replication keeps copies of your data on multiple machines, so the system keeps working if one server fails and can handle more search traffic at once.

Together, they're what let vector databases serve fast search at internet scale, reliably.

Vector Database

Architecture Layers

1

Storage Layer

Holds raw vectors and metadata on disk or in memory.

2

Index Layer

Organizes vectors for fast approximate search (HNSW, IVF, etc.).

3

Query Layer

Handles similarity search, filtering, and re-ranking logic.

4

API Layer

Exposes everything through a simple interface for your app.

How to Evaluate a Vector Search System

Recall

What fraction of the true nearest neighbors did the search actually find?

Precision

Of the results returned, how many were actually relevant?

Latency

How long does a single search take, typically measured in milliseconds?

Throughput (QPS)

How many queries per second the system can handle under load.

Common Pitfalls to Avoid

Bad Chunking

Chunks that are too small or too large hurt retrieval quality badly.

Mismatched Metrics

Using the wrong distance metric for your embedding model.

Ignoring Metadata

Skipping filters leads to irrelevant or stale results.

No Re-ranking

Trusting raw ANN order without a final relevance pass.

Stale Indexes

Not updating vectors when source data changes.

Skipping Evaluation

Shipping without measuring recall or latency first.

Security & Privacy Considerations

Vectors can sometimes be reverse-engineered to reveal fragments of the original data, so they deserve the same protection as the source content.

Key practices: **encrypt data at rest and in transit**, enforce strict access controls per tenant, and avoid embedding sensitive raw data unless truly necessary.

Regulated industries also need to think about data residency and compliance when choosing where vectors are stored.

What Drives Vector DB Costs?

Storage scales with the number of vectors and their dimensionality — more dimensions means more memory per item.

Compute costs come from embedding generation and running searches, especially at high query volume.

Techniques like quantization, smaller embedding models, and caching frequent queries can cut costs significantly without hurting quality much.

Where Vector DBs Are Used Today

RAG for AI Chatbots

Lets AI like ChatGPT pull facts from your own documents before answering.

Recommendation Engines

Finds products, songs, or shows similar to what you already like.

Semantic Search

Search engines that understand intent, not just exact keywords.

Customer Support

Matches incoming tickets to the most similar past resolutions.

More Places Vector Search Shows Up

Fraud Detection

Spots transactions that "look like" known fraud patterns.

Image & Face Search

Finds visually similar images without text tags or labels.

Duplicate Detection

Identifies near-duplicate content even with different wording.

Code Search

Finds functions or snippets with similar logic across a codebase.

Emerging Applications

AI Agent Memory

Gives autonomous agents long-term memory of past interactions.

Drug Discovery

Matches molecule structures with similar chemical properties.

Anomaly Detection

Flags sensor readings or logs that deviate from normal patterns.

Voice & Audio Search

Finds similar-sounding clips, speakers, or music tracks.

Popular Vector Databases

PC

Pinecone

Fully managed, simple to use, popular for production apps

WV

Weaviate

Open source, built-in hybrid search (keyword + vector)

ML

Milvus

Open source, built for massive scale (billions of vectors)

CR

Chroma

Lightweight, beginner-friendly, great for local projects

PG

pgvector

Adds vector search directly inside PostgreSQL

More Vector Databases to Know

QD

Qdrant

Open source, Rust-based, strong filtering performance

VS

Vespa

Battle-tested at scale, built by Yahoo for search and ranking

RD

Redis

Adds vector search to an existing in-memory data store

ES

Elasticsearch

Combines mature full-text search with vector capabilities

FS

FAISS

Meta's library for building your own custom vector index

Choosing the Right Vector Database

If You Need...	Consider
Zero infrastructure to manage	Pinecone
Hybrid keyword + vector search	Weaviate, Elasticsearch
Billions of vectors, self-hosted	Milvus, Qdrant
Quick local prototyping	Chroma
Already using Postgres or Redis	pgvector, Redis

Where Vector DBs Fit in the AI Stack

A modern AI application usually layers several pieces together: an **LLM** for reasoning and generation, an **embedding model** for understanding meaning, and a **vector database** for memory and retrieval.

Orchestration frameworks like LangChain or LlamaIndex tie these pieces together, routing data between the model and the database.

The vector database is often described as the **long-term memory** of an AI system — it's what lets a model "remember" far more than fits in its context window.

Future Trends in Vector Search

Multimodal Embeddings

Single vectors that represent text, images, and audio together.

Serverless Vector DBs

Pay-per-query pricing with zero infrastructure management.

Smarter Agent Memory

Vector stores tailored for long-running autonomous AI agents.

On-Device Vector Search

Lightweight indexes running directly on phones and laptops.

Best Practices Checklist

- Choose a chunking strategy and test it against real queries
- Match your distance metric to your embedding model
- Add metadata filters instead of relying on vectors alone
- Re-rank top results before showing them to users
- Track recall, precision, and latency over time
- Plan for updates — don't let your index go stale

Common Myths

About Vector Databases

Myth	Reality
"They replace SQL databases"	They complement them for meaning-based search
"More dimensions always = better"	More dimensions cost speed and memory too
"ANN search is exact"	It's approximate by design, trading a little accuracy for speed
"Any embedding model works the same"	Model choice hugely affects search quality

Key Terms To Know

Embedding

A list of numbers representing the meaning of a piece of data.

Vector Index

A special structure that organizes vectors for fast similarity search.

ANN Search

Approximate Nearest Neighbor — fast similarity search at scale.

RAG

Retrieval-Augmented Generation — using a vector DB to feed facts to an AI model.

Re-ranking

A precise second pass that reorders top search results by true relevance.

Now You Understand Vector Databases

From embeddings and indexing algorithms to RAG pipelines and production scaling — the technology behind every modern AI search, chatbot, and recommendation system, explained from first principles.



Like this post



Repost it



Follow for more